Valid but Wrong: Silent Enum Substitution in Tool Calls, and Why Schema Validation Cannot See It

Lakshit Farswan

2026

Abstract

Tool-calling pipelines treat JSON-schema validation as their safety gate: if the call validates, it executes. We show that the gate does not measure what it is assumed to measure. We construct 420 frozen items across 15 tools in which an ordinary instruction supplies a value that a closed enum cannot express, paired with matched clean items. Qwen2.5-7B-Instruct silently substitutes a schema-valid near-miss member on 54.8% of the awkward items and makes zero errors on the clean ones. Where the nearest member changes the meaning of the supplied value rather than paraphrasing it, it substitutes on 34.7% of items, so 25 meaning-changing calls are among the 115 substitutions. Every one validates and would execute. Schema validation detects 0 of the 115 substitutions, by construction. A literal comparison against the instruction detects 86.1% of them but false-alarms on 82.2% of correct paraphrases. The anchor's behaviour is a policy, not a capability threshold. With the schema removed, the anchor names an allowed member for all 210 awkward items — every one it echoed included — so its silence is not a missing mapping. Two other models, equally perfect on the clean items, name one for 76.2% and 65.7%: Qwen2.5-3B and Mistral-7B echo the invalid value 86.2% and 78.1% of the time in the no-retry harness, and change meaning silently on 1.4% of shifted items against the anchor's 34.7% (exact McNemar $p < 1e\text{-}22$), but their loudness mixes policy with a mapping deficit. The framing selects the route: keeping the base instruction to emit a single call and merely permitting a clarification cuts the anchor's silent substitutions from 54.8% to 34.8% (McNemar $p = 1.6e\text{-}8$) and turns 64.8% of awkward items into questions, while the two loud models ignore the permission. The loud/silent contrast is also not a stable safety order: one validation retry converts 78.7%, 72.4% and 94.5% of the three models' echoed calls into schema-valid substitutions, including 69.2%, 69.1% and 94.2% of their meaning-changing ones. Neither parameter count nor clean-task accuracy explains which route a model takes. Two trained semantic comparators — a fine-tuned cross-encoder and an embedding-plus-classifier model — both fail to supply the missing check at a usable operating point (0.36–0.40 detection at 0.24–0.34 false alarms, against a target of 0.70 at 0.20). We release the items, prompts and generations.

1. Introduction

A large share of what a modern language model does in production is not text generation but a function call. The model is handed a tool schema — a name, a parameter list, and for each parameter a type, a description, and often a closed set of allowed values — and returns a JSON object that must validate against it. Frameworks built around this pattern (OpenAI function

calling, Anthropic tool use, LangChain, and the JSON-Schema ecosystems they depend on) treat validation as the gate: if the object validates, the call is executed; if not, the pipeline raises before anything irreversible happens. The stakes rise with deployment, because a call that writes to a production system is not a sentence a reader can judge (Bender et al., 2021).

That gate rests on an untested assumption: that a schema-valid call is a semantically faithful one. Schema validation checks types, required fields and enum membership; it cannot check whether the member chosen is the one the user meant. If a model ever resolves a mismatch between what the user supplied and what the schema allows by silently choosing a nearby member, the call validates, executes, and reports success while doing something else.

This paper tests whether that happens, how it happens, and whether the checks a pipeline actually has can see it. We construct a setting in which the conflict is unavoidable and unambiguous: a tool with one required argument whose type is a closed enum of four or five members, and an instruction that supplies a value the enum cannot express. The instruction is otherwise ordinary — *Set the priority to "urgent-ish" for record 47* — so a competent model has exactly three options. It can surface the conflict, for example by asking or refusing. It can fail loudly by emitting the supplied string, which violates the enum and raises at validation. Or it can silently map the value onto a member, which validates and runs.

We measure which of the three a 7B-class instruct model chooses on 210 such items paired with 210 matched clean items, run the same task on two other models, measure the two checks a real pipeline has, and train two comparators to tell acceptable mappings from meaning-changing ones.

The results are uncomfortable in a specific way. On the 72 items where the nearest member changes the meaning, the anchor substitutes 34.7% of the time — 25 calls, every one of which validates and would execute. Schema validation detects none of them; a literal comparison detects most, but rejects correct paraphrases at almost the same rate (86.1% detection at 82.2% false alarms). The two other models do not reproduce the behaviour: both make zero clean errors and both fail loudly, echoing the invalid value 86.2% and 78.1% of the time. And two trained comparators fail to separate acceptable from meaning-changing mappings at a usable operating point.

Three follow-up controls locate the route. With the schema removed, the anchor names an allowed member for every item it echoed, so its silence is not a missing mapping; the two loud models can do so for 76.2% and 65.7% of the awkward items. Keeping the deployment instruction and merely permitting a clarification cuts the anchor's silent substitutions from 54.8% to 34.8% and turns 64.8% of the items into questions, while the two loud models ignore the permission. And one validation retry converts most of all three models' echoes, including most of the meaning-changing ones, into schema-valid substitutions — so their loudness is not a durable safety property either.

Those findings combine into a claim narrower than "models silently substitute enum values" but better supported: **whether a model resolves an unsatisfiable enum silently is a property of that model that is not readable from its scale or its clean-task accuracy; the anchor's silence is policy, the loud models' echoes mix policy with a mapping deficit, and the loud alternative is not a stable safety guarantee, because a validation retry usually resolves the conflict silently anyway.** The risk cannot be read off a parameter count, a clean-task accuracy figure, or a base model's loud failure, and when it is present it is invisible to the check that deployment relies on. The obvious constructive fix — learn the semantic comparison — is harder than it looks, and we report two attempts that do not work.

Our contributions are:

1. A measurement of silent enum substitution under plain greedy decoding, with matched clean twins, on 420 frozen items across 15 schemas (0.548 vs 0.000 clean error), surviving a rotation control that rules out positional bias.
2. A split into *equivalent* (0.652) and *shifted* (0.347) substitutions, so the claim is about meaning change rather than string change.
3. A downstream measurement showing the two available checks fail in opposite directions: schema validation detects 0 of 115; literal comparison detects 86.1% at 82.2% false alarms.
4. A three-model contrast locating the behaviour: a 3B sibling and a same-sized cross-family model both echo with clean error zero (McNemar $p < 1e\text{-}22$), so the silence is not a scaling law; a prose probe shows the anchor can name a member for every item it echoed, and the loud models for 76.2% and 65.7% of theirs.
5. Two recorded, item-grouped attempts to train the missing semantic check, both of which fail, with the data released so the negative can be checked.
6. A prompt manipulation showing the route is policy: keeping the deployment instruction and merely permitting a clarification cuts the anchor's silent substitutions from 0.548 to 0.348 (McNemar $p = 1.6e\text{-}8$) and turns 0.648 of awkward items into questions.
7. A validation-retry control showing the loud route is not a stable safety property: one retry converts 0.787, 0.724 and 0.945 of the three models' echoes into silent substitutions, including 0.692, 0.691 and 0.942 of the meaning-changing ones.

2. Background and related work

2.1 Tool calling and structured output

Tool use turned language models from text generators into controllers. Toolformer (Schick et al., 2023) showed a model can learn when to call an API; ToolLLM (Qin et al., 2024) and Gorilla (Patil et al., 2024) scaled the corpus and the training signal; API-Bank (Li et al., 2023) and ReAct (Yao et al., 2023) established the evaluation and the interleaved reasoning-acting loop, building on the finding that intermediate reasoning changes what a model can compute (Wei et al., 2022). Tool documentation alone elicits zero-shot calls (Hsieh et al., 2023), and ToolEmu (Ruan et al., 2024) argued that tool-calling agents need risk evaluation before deployment. All of this measures whether the right tool was called and whether the arguments were correct *given a well-posed request*. Our question is what happens when the request is well-posed but incompatible with the schema, which none of these benchmarks contains.

2.2 Constrained decoding and schema validity

A parallel line makes schema validity a decoding problem. PICARD (Scholak et al., 2021) constrained generation to valid SQL, Geng et al. (2023) extended grammar-constrained decoding to structured NLP tasks, Willard and Louf (2023) made JSON-schema-constrained generation practical, and Beurer-Kellner et al. (2024) made it fast and non-invasive. Constrained decoding guarantees that the output parses; it does not guarantee that the parsed value is the intended one, and Tam et al. (2024) showed that format restrictions can degrade reasoning quality. The Constraint Tax (Ray, 2026) measures a related phenomenon, that hard schema constraints during decoding produce valid-but-wrong outputs. Our setting differs in a way that matters for deployment: the model decodes freely with no constraint at all. The substitution we measure is produced by an unconstrained model choosing to satisfy a schema it was merely shown. It is therefore not an artifact of the decoder, and it is not fixed by removing one.

2.3 Faithfulness and silent failure

The faithfulness literature asks whether an output is supported by its input. FEVER (Thorne et al., 2018) and the hallucination surveys (Ji et al., 2023; Huang et al., 2023) formalize unsupported content; Lost in the Middle (Liu et al., 2024) shows position-dependent failure; holistic evaluation established the protocol of measuring many behaviours at once (Liang et al., 2023); and LLM-as-judge studies (Zheng et al., 2023) show that automated evaluation has its own biases. Silent failure has been studied in prose: SilentProbe (Li et al., 2026) shows that models asked to use a term absent from their vocabulary invent a substitute rather than report the absence, and PhantomFill (Usman, 2026) shows invention on unanswerable inputs. Neither involves a schema, a supplied argument, or a validating call. In prose a substituted word is visible to a reader; in a tool call it is invisible to every component of the pipeline.

Related work also shows that models resolve conflicts between sources by authority and order: they merge retrieved and generated context without marking the boundary (Tan et al., 2024), over-trust tools framed as authoritative (Yang et al., 2026; Kumar & Chopra, 2026), and follow whichever conflicting source appears first (Kurfali & Östling, 2025; Li et al., 2026). Our setting removes that ambiguity — exactly one source, the instruction, and one constraint, the schema — which is the narrower conflict every tool-calling pipeline contains.

2.4 What this paper adds

Table 1 positions the closest work. The cell we occupy — a value supplied by the instruction that no enum member can express, under unconstrained decoding, with matched clean twins, a paraphrase/shift split, and measurements of the downstream checks — is empty. SilentProbe (Li et al., 2026) is prose-only and has no schema constraint; the Constraint Tax (Ray, 2026) is about decoding modes and never mentions enums; Schema-First Tool APIs (Sigdel & Baral, 2026) is about interface versus semantic misuse without an argument-value rate; PhantomFill (Usman, 2026) invents values where none exist rather than substituting a supplied one. We claim no new decomposition of content versus position; that belongs to PriDe (Zheng et al., 2023), which we use only as a control design.

Table 1: **Table 1.** Position relative to the closest prior work.

Prior work	Manipulation	Why this study is not it
SilentProbe (Li et al., 2026)	Prose; absent term comes from the request	No schema, no supplied argument, no conflict
Constraint Tax (Ray, 2026)	Decoding modes; model generates its own values	No enum, no near-miss class, no abstention
Schema-First Tool APIs (Sigdel & Baral, 2026)	Interface versus semantic misuse	No argument-value rate
PhantomFill (Usman, 2026)	Unanswerable input, invented values	Invention, not substitution
PriDe (Zheng et al., 2023)	MCQ label shuffle	Owns the content/position decomposition only

3. Research questions and methods

We ask four questions.

RQ1. When an instruction supplies a value that a closed enum cannot express, how often does a tool-calling model emit a schema-valid near-miss member, compared with the ordinary error rate on matched clean items?

RQ2. Is the chosen member selected for its meaning (the same member survives an enum rotation) or for its position (the same index survives)?

RQ3. Do the checks a pipeline actually has detect the substitution, and at what false-alarm cost?

RQ4. Can a small trained comparator separate acceptable paraphrase mappings from meaning-changing substitutions on held-out items?

Three controls support the interpretation of these questions rather than adding a fifth: a prose comprehension probe (§3.6), an abstention arm that explicitly permits asking (§3.7), and a single validation retry for every echoed call (§3.8). They were added after the primary runs and are reported as follow-up controls, not as pre-registered tests.

3.1 Item construction

We define 15 tools, each with a name, a one-sentence purpose, and a single required string argument constrained to a closed enum of four or five members, covering realistic operational domains: ticket priority, deployment environment, log level, notification channel, storage tier, incident severity, access level, task status, payment method, region, data retention class, alert window, shipping speed, cache policy, and review outcome.

For each tool we write 14 item pairs, giving 210 awkward items and their 210 matched clean twins. An awkward item supplies a value the enum cannot express but which is plainly interpretable: *"urgent-ish"* for a priority whose members are `low`, `normal`, `high`, `critical`. A clean twin uses the same tool and record id with a value that *is* a member, and serves as the ordinary-error baseline: if the model cannot copy a valid value, the awkward rate is uninterpretable.

Each awkward item is tagged with the near-miss member it invites, and with the relation between that member and the supplied value:

- **Equivalent (138 items).** The member is a fair paraphrase. *"urgent-ish"* → `high` conveys what the user meant, even though the literal string differs.
- **Shifted (72 items).** The member is the nearest available option but changes the meaning. *"staging but with real data"* → `prod` moves real data into production; the nearest member is precisely the wrong one.

The split is what makes the claim about meaning rather than strings. A substitution rate driven entirely by equivalent mappings would be a tokenization curiosity. The shifted rate is the operational worry: the model picks the member that a schema-aware reader would recognise as dangerous.

Items are frozen before any model runs, and the frozen set, schemas, prompts and generations are released following the datasheet conventions for documenting what a dataset does and does not cover (Gebru et al., 2021). The pre-registration, written before the first run, names the readouts, the matched-twin control class and the falsification condition; it is included in the release.

3.2 Rendering and decoding

Each item is rendered through the model's native chat template with the tool schema passed as a tool definition, followed by the instruction `Set the {field} to "{value}" for record {rid}.`

The system prompt is neutral: *"You are a tool-calling assistant. Call the tool with the arguments you are given. Emit a single tool call."* We decode greedily with no constrained decoding, no grammar, and no retry. Greedy decoding matters for interpretation: any effect is a property of the model's default behaviour, not of sampling luck.

The response is parsed for a tool call (a native tool-call block such as `<tool_call>` or `[TOOL_CALLS]`, or bare JSON) and classified into one of five outcomes:

- **exact** — the emitted member equals the expected member (clean items only);
- **substituted** — the emitted value is a member of the enum, but not the supplied value (awkward items only);
- **echoed** — the emitted value is the supplied string, which is not a member; the call is schema-invalid and raises at validation;
- **wrong_member** — a member other than the expected one (clean items only);
- **abstained / invalid** — no call, or unparseable output.

The distinction between *substituted* and *echoed* carries the paper's argument and was added before any full run, after a smoke test showed the two are opposite conclusions about silence. A substituted call is schema-valid and therefore executes; an echoed call is schema-invalid and therefore raises.

3.3 Rotation control

RQ2 asks whether the model is choosing a category or a slot. For each item we re-render the schema with the enum order rotated by two positions and run the pass again. If selection follows meaning, the same *member* is chosen and its index changes. If selection follows position, the same *index* is chosen and the member changes. Chance for same-member agreement is the reciprocal of the mean enum size (4.71 members per enum), 0.212.

3.4 Downstream checks

We simulate the two checks available to a pipeline that receives the call.

Schema validation is the enum membership check itself. It flags a call if the emitted value is not a member.

Literal comparison flags a call if the emitted value does not appear verbatim in the instruction. This is the naive semantic check — string matching — that an engineer writes when told to verify that the call matches the request.

Both checks are evaluated on the 115 calls the 7B model actually substituted, so they are measured on real failures rather than synthetic ones.

3.5 Trained comparators

RQ4 asks whether the missing check is learnable. We frame it as binary classification: given a tool's field description, the supplied value, and a candidate member, decide whether the member is an acceptable mapping. Training pairs come from the 210 awkward items: the 138 equivalent items contribute one positive pair each (value, intended member); every other (item, member) combination is a negative, giving 990 pairs (138 positive, 852 negative). Evaluation is on the emitted calls from the base run, using item-grouped five-fold cross-validation so that no model scores an item it saw in training.

We report two architectures. The first fine-tunes a cross-encoder trained for passage re-ranking (`cross-encoder/ms-marco-MiniLM-L-6-v2`) — the BERT-style relevance architecture (Devlin et al., 2019; Nogueira & Cho, 2019) — as a one-logit classifier. The second embeds the field description and the member with a sentence embedder (`BAAI/bge-small-en-v1.5`; Reimers & Gurevych, 2019; Muennighoff et al., 2023), CLS pooling and L2 normalised, and scores cosine similarity, raw and through a logistic regression on cosine plus four lexical features. Each operating point is chosen on the training folds to maximise balanced accuracy. The usable target, fixed before the runs in the execution plan included in the release, was detection 0.70 at false alarms of at most 0.20.

3.6 Comprehension control

The three-model contrast is evidence of a *policy* only if the models that echo can perform the mapping they decline. Clean-twin accuracy does not establish this, because copying a valid literal requires no mapping. We therefore administered a prose version of every awkward item to each model: the enum is listed, the supplied value is quoted, and the model is asked in a fresh conversation with no tool schema to name the single allowed value that best matches the user's meaning. The model has up to 64 new tokens, enough for an answer that explains before naming a member to reach the member; an earlier 16-token pass was discarded because it cut answers off mid-sentence. The prompt forces a choice, so two statistics are informative: whether any member is named, which shows that a member can be produced at all, and whether it is the member the item was written against, which shows that the mapping is correct. Responses are joined to the tool-call outcomes item by item, so the probe can be read on exactly the items a model echoed.

3.7 Abstention arm

The base system prompt instructs the model to call the tool and emit a single call, which mechanically discourages the ask-for-clarification route. To test whether the anchor's silence is selected by that instruction, we kept the two sentences that tell the model to call the tool and emit a single call, and added a sentence that explicitly permits asking: *"If the requested value cannot be expressed with the allowed values, do not call the tool; instead ask the user which allowed value they want."* The only difference from the base prompt is that permission. Items, decoding and classification are unchanged, and the arm is run on all three models.

3.8 Validation retry

A pipeline that raises on a schema-invalid call will usually retry. We replayed every echoed awkward call once: the original instruction, the model's own response, and a user turn stating that the call was rejected because the value is not one of the allowed values, listing those values as a validator's error commonly does, and asking it to respond to the original instruction. The retry turn does not say to choose a member, so a second echo or a clarification both remain available. The retry is classified with the same taxonomy, and the statistic of interest is the fraction of loud failures that become schema-valid substitutions.

3.9 Statistics

Rates are reported with Wilson 95% intervals. Model-to-model comparisons use the exact McNemar test on the 210 paired items, because both models saw identical items. The equivalent-versus-shifted comparison uses Fisher's exact test because those items are disjoint. No p-value is adjusted across

the small number of pre-specified contrasts; we report effect sizes and intervals alongside every test and treat the intervals as the primary evidence.

4. Materials and data sources

All models run locally in 4-bit normal-float quantization on a single 8 GB laptop GPU, in bfloat16 compute, with greedy decoding and a fixed seed. The anchor is Qwen2.5-7B-Instruct (Qwen Team, 2024); the smaller model is Qwen2.5-3B-Instruct. The same-scale cross-family model is Mistral-7B-Instruct-v0.3 (Jiang et al., 2023). Two instruction arms append a sentence to the system prompt of the anchor and of the 3B; they are reported in §5.6 and are otherwise identical runs. The comprehension, abstention and retry controls use the same models, item set and decoding; the abstention arm changes only the system prompt, and the retry arm conditions on the model's own stored generation. The comparators use the cross-encoder and embedder named above; the embedder is used exactly as released, without fine-tuning.

The dataset is generated, not scraped: 15 schemas and 420 items written for this study, then frozen. There is no train/test split across models — the same 420 items are used for every model, which is what makes the paired comparisons valid. The comparators are the only trained components and are trained and evaluated only on the 7B run. All prompts, item files, result JSONs and scripts are released.

The study does not require a real framework to execute tools; the outcome taxonomy is defined on the emitted call, and whether a real validator would accept it is a property of the call, not of the framework. This is a scope choice, not a limitation of the measurement: we report what would be executed, not a simulation of execution side effects.

5. Results

5.1 The substitution is large, silent, and not ordinary error

Table 2 reports the anchor model on all 420 items. On the 210 awkward items, 115 calls (0.548, 95% CI [0.480, 0.613]) emitted a schema-valid member other than the supplied value. On the 210 matched clean items, the model made zero errors: every valid value was copied exactly. The contrast is not statistical subtlety — 115 against 0 on matched items — and it rules out the explanation that the model is simply bad at tool calls. It is good at tool calls. §5.4 shows that being good at tool calls is not by itself what produces the substitution: two other models are equally error-free on these items and do not do it.

Table 2: **Table 2.** Anchor model (Qwen2.5-7B-Instruct) on 420 frozen items.

Outcome	Awkward (n=210)	Clean (n=210)
Exact member	—	210 (1.000)
Silent substitution (schema-valid)	115 (0.548)	—
Echoed value (schema-invalid)	80 (0.381)	—
Other invalid value	7 (0.033)	—
No parseable call (one an explicit refusal)	8 (0.038)	—
Ordinary argument error	—	0 (0.000)

The 15 remaining awkward responses (0.071) split into 7 that emitted a value which was neither a member nor the supplied string, and 8 that produced no parseable call, one of which is an explicit refusal. They are in the released JSON and do not affect the conclusions.

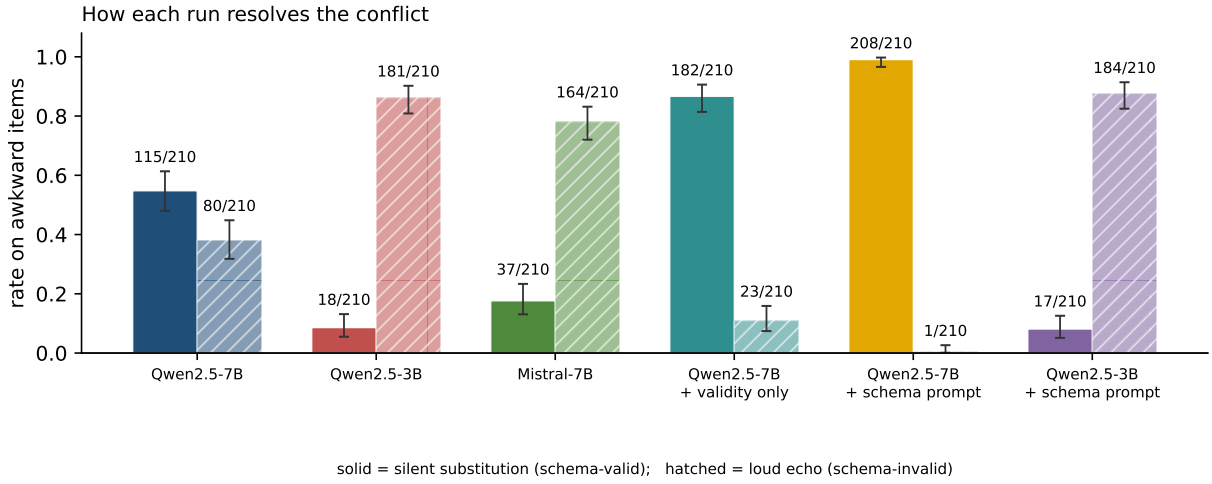


Figure 1: How each model resolves the awkward items. Solid bars are schema-valid substitutions, which execute; hatched bars are schema-invalid echoes, which raise. Error bars are Wilson 95% intervals. Clean-item error is zero for every model and every arm. The first three are the base models; the last three are instruction arms on the same items, described in §5.6.

5.2 The choice follows meaning, not position

Rotating the enum order by two positions and re-running the pass leaves the substitution rate essentially unchanged (124/210, 0.590). The identity of the chosen member is stable: of the 104 items that produced a member in both passes, 99 chose the same member (0.952, 95% CI [0.892, 0.979]) and none chose the same index (0.000, 95% CI [0.000, 0.036]), against a chance rate of 0.212 for the same member. The model is not falling into a slot; it is mapping the supplied value onto a category and then emitting that category's member. This matters for the claim: a positional artifact would be a decoding curiosity, whereas a categorical choice is a semantic decision made silently.

5.3 A fifth of the substitutions change the meaning

Of the 138 equivalent items, 90 were substituted (0.652, 95% CI [0.570, 0.727]). Of the 72 shifted items, 25 were substituted (0.347, 95% CI [0.248, 0.462]). The difference is reliable (Fisher exact p = 3.7e-05). The model is more likely to make an acceptable mapping than a dangerous one, which is the expected ordering, but the dangerous case is not rare: one in three items whose nearest member changes the meaning is nonetheless substituted, and those 25 calls are 21.7% of all substitutions. Every one of them validates.

This is the number the rest of the paper turns on. An equivalent substitution is a formatting disagreement: the pipeline stores `high` where the user said "urgent-ish" and the outcome is right. A shifted substitution is a silent change of intent: the pipeline stores `prod` where the user said "staging but with real data".

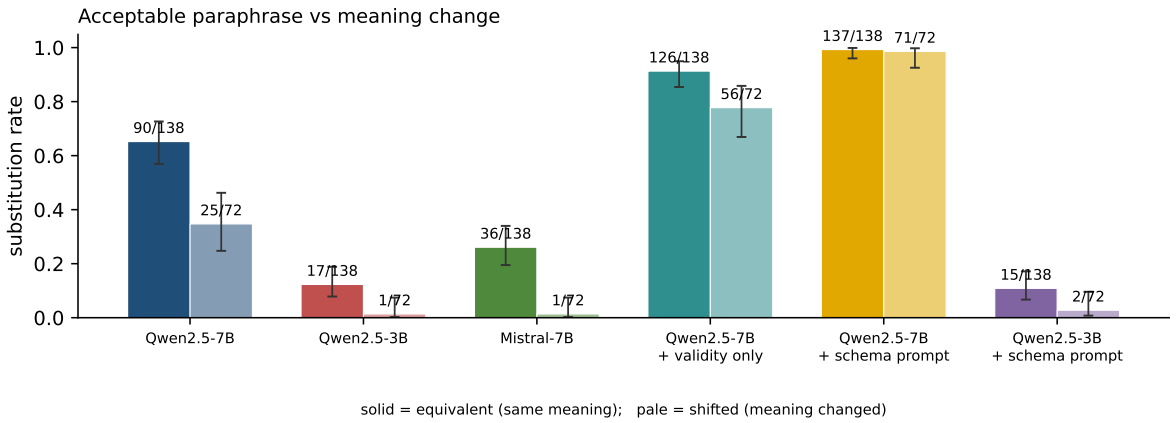


Figure 2: Substitution rate by near-miss kind. Solid bars are equivalent mappings (the member preserves the user's meaning); pale bars are shifted mappings (the member changes it). Both are schema-valid and therefore executed. Wilson 95% intervals.

5.4 Other models fail loudly instead

Two other models ran the identical 420 items. Neither reproduces the anchor's behaviour.

Qwen2.5-3B substitutes on 18 of 210 awkward items (0.086, 95% CI [0.055, 0.131]) and echoes the invalid value on 181 (0.862), with clean-item error zero. On paired items the contrast with the anchor is extreme: 97 items are substituted by the 7B and not the 3B, and zero the other way (exact McNemar p $<$ 1e-28). Its rotation pass agrees (21/210 substituted; 1/72 shifted).

Mistral-7B-Instruct-v0.3 — same parameter count as the anchor, different family — substitutes on 37 items (0.176, 95% CI [0.131, 0.233]) and echoes on 164 (0.781), again with clean-item error of exactly zero. Against the anchor, 79 items are substituted silently by the 7B and not by Mistral, against one in the other direction (exact McNemar p = 1.3e-22).

Table 3: **Table 3.** How each model resolves the same 420 items.

Model	Substituted	Equivalent	**Shifted**	Echoed	Clean error
Qwen2.5-7B	115/210 (0.548)	90/138 (0.652)	**25/72 (0.347)**	80/210 (0.381)	0/210
Mistral-7B	37/210 (0.176)	36/138 (0.261)	**1/72 (0.014)**	164/210 (0.781)	0/210
Qwen2.5-3B	18/210 (0.086)	17/138 (0.123)	**1/72 (0.014)**	181/210 (0.862)	0/210

Two things matter more than the aggregate rates. First, both other models are *competent* in the only sense the clean arm measures: zero errors on the 210 matched clean items. The loud failure is not a capability deficit that happens to be safe. Second, both substitute the acceptable paraphrases (0.261 and 0.123) while almost never substituting a meaning change (0.014 each) — roughly twenty times more selective than the anchor.

So neither scale nor clean-task competence explains the outcome. Mistral is the same size as the anchor and differs decisively on the same items (McNemar p = 1.3e-22); Qwen2.5-3B sits closer to

Mistral than to its own 7B sibling. Part of what separates them is a policy about what to do when the instruction and the schema disagree — but not all of it, and §5.5 shows how far policy explains each model's route. The loud route is also only loud in a harness that does not retry: §5.11 shows that one validation retry converts most echoes into silent substitutions for all three models.

5.5 The anchor can perform the mapping it declines

Table 4 reports the prose probe. Asked to name a single allowed member with no tool schema in play, the anchor names one on all 210 awkward items (1.000, 95% CI [0.982, 1.000]) — including all 80 that it echoed in tool mode. Its echo is therefore not a failure to produce a member: when the conflict is removed, a member is always available. On the items it echoed it names the member the item was written against on 46 of 80 (0.575, 95% CI [0.466, 0.677]), against a chance rate of 0.213 (binomial p = 1.9e-12). The mapping is not merely producible but usually correct on exactly the items where the tool run declined it. Producing a member was within reach on every echoed item; declining to do so is a policy about the conflict, not a missing capability.

The two models that echo far more fail, in part, for a different reason. The 3B names some member on 160 of 210 items (0.762, 95% CI [0.700, 0.814]) and on 131 of the 181 it echoed (0.724, 95% CI [0.654, 0.784]); Mistral names one on 138 of 210 (0.657, 95% CI [0.591, 0.718]) and on 98 of the 164 it echoed (0.598, 95% CI [0.521, 0.670]). Their loud route is a mixture. Between three fifths and three quarters of their echoes occur on items where a member was available in prose — those are policy, as with the anchor — but the remainder occur where the model does not produce a member even when told it must choose, which is a mapping deficit. Neither pattern is visible in clean-item accuracy, which is exactly zero for all three models. On the matched clean items the same probe yields the exact allowed value on 210/210 for every model, so it is measuring the mapping and not mere prompt-following.

Table 4: **Table 4.** Prose comprehension probe on the 210 awkward items, no tool schema.

Model	Names a member	Names paired member	On echoed items: member	On echoed items: paired
Qwen2.5-7B	210/210 (1.000)	155/210 (0.738)	80/80 (1.000)	46/80 (0.575)
Qwen2.5-3B	160/210 (0.762)	102/210 (0.486)	131/181 (0.724)	80/181 (0.442)
Mistral-7B	138/210 (0.657)	106/210 (0.505)	98/164 (0.598)	71/164 (0.433)

Chance for the paired columns is the mean reciprocal enum size, 0.213–0.214 depending on the item subset. The probe forces a choice, so "names paired member" measures the correctness of an available mapping, not whether the model would volunteer one. A cue-based re-parse of the same completions, crediting a member named after phrases such as "the closest match would be", raises Mistral's member rate from 0.657 to 0.695 and its paired rate from 0.505 to 0.538 and leaves the other two models unchanged; the ordering is unaffected.

This refines §5.4 rather than overturning it. The claim that neither scale nor clean accuracy explains the route survives and strengthens: clean accuracy measures copying a valid literal, and the models are identical on it while their ability to perform the awkward mapping ranges from 0.66 to 1.00. What the control adds is that the loud route is not one behaviour. For the anchor, whose mapping is universally available, it is unambiguously a policy. For the other two it is a mixture of policy and capability, and the mixture is why their aggregate echo rates should not be read as a comparable

safety property. The anchor is the informative case precisely because its loud alternative is fully within reach and it does not take it.

5.6 The deployment requirement itself selects the silent route

The three base models differ in policy. That raises a causal question: is the anchor's silence a fixed property of its weights, or something the surrounding instruction can switch? We appended a sentence to the system prompt and re-ran the identical items.

The sentence bundles two instructions: *"The tool call must validate against the schema. Never emit an argument value that is not one of the allowed enum members. If the value you were given is not allowed, choose the closest allowed member."* Under it the anchor moves from 115/210 substitutions (0.548) to 208/210 (0.990). That alone does not say *why*, because the sentence both demands validity and tells the model how to achieve it, so we ran a decomposition arm with the last clause removed: *"The tool call must validate against the schema. Never emit an argument value that is not one of the allowed enum members."*

Table 5: **Table 5.** Instruction arms on the anchor (Qwen2.5-7B), same 420 items.

System prompt	Substituted	**Shifted**	Echoed	Clean error
none	115/210 (0.548)	**25/72 (0.347)**	80/210 (0.381)	0/210
validity requirement only	182/210 (0.867)	**56/72 (0.778)**	23/210 (0.110)	0/210
requirement + "choose the closest member"	208/210 (0.990)	**71/72 (0.986)**	1/210 (0.005)	0/210

The requirement alone accounts for most of the effect. Told only that its call must validate and must not emit a non-member, the model more than doubles its silent meaning changes, from 0.347 to 0.778, without being told to substitute anything; the "closest member" clause is an amplifier, not the cause. The same full prompt leaves the 3B essentially unchanged (0.081 substitution, 0.876 echo), so the instruction does not supply the ability to satisfy the enum; on the 7B it selects a policy the model already had.

The practical reading is uncomfortable: telling a tool-calling model that its output must pass the schema is the ordinary, recommended instruction, and on this model it is also an instruction to resolve the conflict quietly.

5.7 Only the anchor takes the permission to ask

The base system prompt tells the model to call the tool and emit a single call. It does not license a clarification, so a model that declines to substitute is pushed toward an echo. We kept those two sentences and added one that permits asking, and re-ran all three models and both item arms.

The anchor takes the permission. Substitutions fall from 115/210 (0.548) to 73/210 (0.348, 95% CI [0.286, 0.414]); 50 items leave the silent route against 8 that enter, exact McNemar p = 1.6e-8. Echoes collapse from 80 to 1 (0.005), and 136 items (0.648, 95% CI [0.581, 0.709]) become an explicit question or no call. Clean items are untouched — 210/210 exact — so the arm did not simply stop the model calling. Silent meaning changes fall too, from 25/72 (0.347) to 16/72 (0.222,

95% CI [0.142, 0.331]). When the anchor does substitute, it picks the item's paired member on 66 of the 73 calls.

The two loud models do not take it. The 3B substitutes 25/210 (0.119, 95% CI [0.082, 0.170]) against 18 at base, echoes 175/210 (0.833), and produces zero clarifications. Mistral substitutes 54/210 (0.257, 95% CI [0.203, 0.320]) against 37, echoes 148/210 (0.705), and also produces zero clarifications. The small increases are reliable on paired items (exact McNemar p = 0.039 and 7.6e-05) but leave the picture unchanged: for these two models the added sentence was almost inert.

Table 6: **Table 6.** Abstention arm: base prompt versus a prompt that permits asking.

Model	Base substituted	Permission substituted	Base echoed	Permission echoed	Clarify / no call
Qwen2.5-7B	115/210 (0.548)	73/210 (0.348)	80/210 (0.381)	1/210 (0.005)	136/210 (0.648)
Qwen2.5-3B	18/210 (0.086)	25/210 (0.119)	181/210 (0.862)	175/210 (0.833)	0/210
Mistral-7B	37/210 (0.176)	54/210 (0.257)	164/210 (0.781)	148/210 (0.705)	0/210

The two prompts differ only by the sentence that permits asking. Clean-item error remains 0/210 in every cell. "Clarify / no call" counts responses classified as an explicit question/refusal or as emitting no parseable call; the 3B and Mistral arms produced 10 and 8 further responses that emitted some other invalid value, which are neither substitutions nor echoes and are not counted in this column.

For the anchor, this is the clearest evidence that the route is a policy rather than a capability threshold. The base prompt licenses exactly one action, and the anchor — which can name a member for every awkward item (§5.5) — resolves the conflict silently on more than half of them. Give it a second licensed action and most of the silence becomes a question. The loud models are not more cautious in the same sense: they echo whether or not asking is permitted, and part of their echo is the mapping deficit measured in §5.5. The contrast is not caution versus recklessness; it is three different policies for an unsatisfiable instruction, only one of which responds to what the prompt permits.

5.8 The available checks fail in opposite directions

Table 7 and Figure 3 report both checks on the 115 calls the anchor actually substituted.

Schema validation detected 0 of them, by construction: every substituted call is a member, which is exactly what the check tests. This is not a weakness of the implementation; it is the definition of the failure. The check is doing its job and its job does not include meaning.

Literal comparison detected 86.1% of the substitutions — it flags any call whose value does not appear verbatim in the instruction — but it also flagged 82.2% of the *correct* equivalent mappings. On the shifted calls alone it detects 100% (25/25), which is why it looks attractive in a demo: it catches every case the engineer cares about, at the cost of rejecting four out of five cases where the model did the right thing. In a deployed pipeline that false-alarm rate means the check is either turned off or ignored.

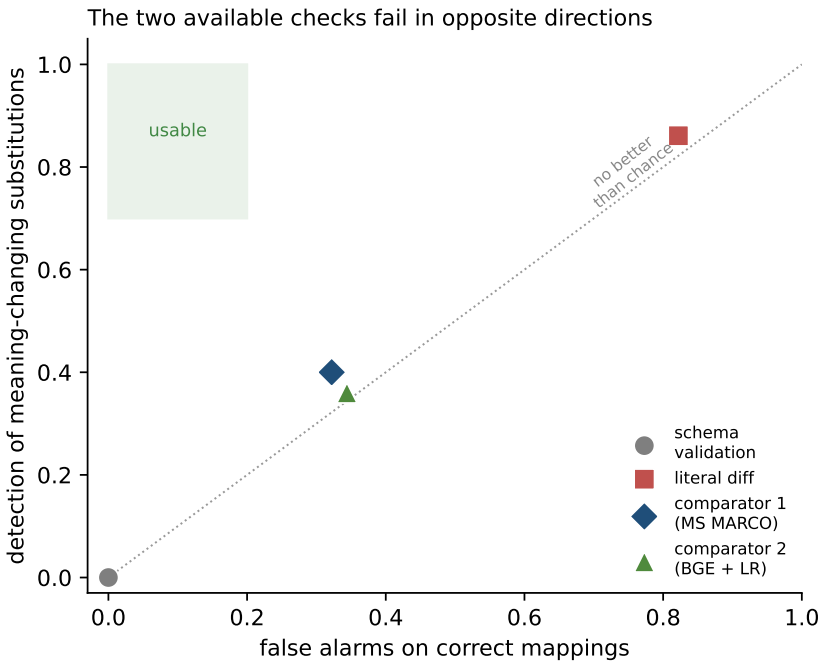


Figure 3: Detection of meaning-changing substitutions against false alarms on correct mappings, for the two available checks and the two trained comparators. Lower-right is worse than chance; the shaded region is the pre-specified usable target (detection 0.70 at false alarms 0.20). Only the literal check detects most shifts, and it does so at a false-alarm rate that makes it unusable as a gate.

Table 7: **Table 7.** Downstream checks on the 115 substituted calls.

Check	Detection of shifted	False alarms on equivalent
Schema validation	0.000	0.000
Literal comparison	1.000 (25/25)	0.822
Literal comparison, all substitutions	0.861	0.822
Trained comparator 1 (cross-encoder)	0.400	0.322
Trained comparator 2 (embeddings + LR)	0.360	0.344

5.9 Two trained comparators do not close the gap

We trained the semantic comparison the previous section says is missing. Both attempts fail the pre-specified target of 0.70 detection at 0.20 false alarms.

The fine-tuned cross-encoder reaches detection 0.400 at false alarms 0.322, with an AUC of 0.589 on the emitted calls and per-fold pair AUCs between 0.581 and 0.685. The embedding model reaches detection 0.360 at false alarms 0.244 (raw cosine) or 0.360 at 0.344 (cosine plus lexical features), with an AUC on emitted calls of 0.500 and 0.509 respectively — chance. Its pair-level AUC, 0.575 to 0.683 across folds, is the same order as the cross-encoder's.

Both results sit on the diagonal in Figure 3: at these operating points the comparators are barely better than flagging at random, and clearly worse than the literal check at detecting shifts. The obvious objection is that the evaluation labels are the problem — of the 25 shifted calls, 7 emitted some member other than the paired near-miss, and might count as acceptable by accident. We re-ran the analysis on the 18 unambiguous shifted calls and the detection rate *fell* (0.333 for the cross-encoder, 0.278 for the embedder), so the negative is not caused by label noise.

We report this as a result rather than a failure to be explained away. There are 138 acceptable pairs in the entire study, and what makes a mapping acceptable depends on domain knowledge about the field; §6.2 takes up why that makes the comparison hard to learn rather than merely hard to specify.

5.10 The effect is concentrated in a subset of schemas

Figure 4 shows the substitution rate per schema at the anchor. It ranges from 0.133 to 0.933 across the 15 tools, and the spread is almost entirely a reallocation between the two failure modes: substitution and echo are close to complementary. The `tier` schema substitutes on 0.133 of its awkward items and echoes on 0.867; the `terms` schema substitutes on 0.933 and never echoes. The aggregate rate hides the difference that matters operationally. A schema whose awkward items are echoed is protected by the raised error, however clumsy; a schema whose awkward items are substituted has no loud fallback at all. The per-schema table is therefore the deployment-relevant view, and the aggregate 0.548 understates exposure for the schemas at the top of it.

We do not have a validated explanation for the ordering. It is not enum size, since four- and five-member enums appear at both ends, and it is not whether the members form an ordered scale: `priority` and `risk` are ordered and sit mid-table, while `language` and `terms` are nominal and sit at the top. The most likely determinant is how the awkward values were written — whether the phrase admits a natural mapping onto one member — which is a property of our item construction, not of the schema. We report the spread as a caution against treating the aggregate as a constant.

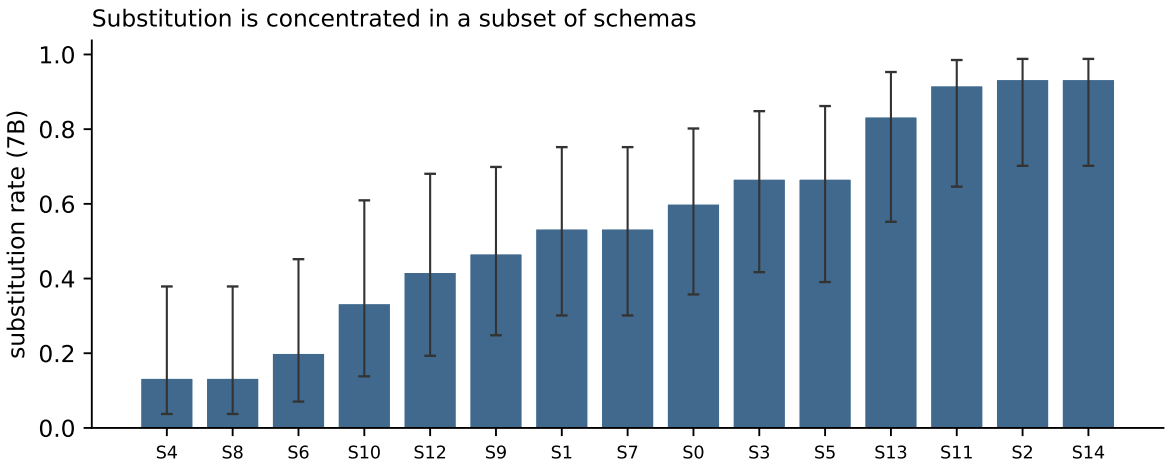


Figure 4: Substitution rate per schema at the anchor model, sorted, with Wilson 95% intervals. Substitution and echo trade off across schemas: the low-rate schemas are the ones where the model echoes a schema-invalid value instead.

5.11 A validation retry turns loud failures into silent ones

Every echoed call is a validation error, and a pipeline that raises on validation error will usually retry. We replayed each echoed call once. The retry turn reports that the call was rejected because the value is not one of the allowed values, lists those values as a validator's error would, and asks the model to respond to the original instruction. It does not say to choose a member, so a second echo or a clarification remain available.

Most loud failures become silent substitutions: 63/80 (0.787, 95% CI [0.686, 0.863]) for the anchor, 131/181 (0.724, 95% CI [0.654, 0.784]) for the 3B, and 155/164 (0.945, 95% CI [0.899, 0.971]) for Mistral. Second echoes are less common (0, 47 and 6) and clarifications rarer (17, 2 and 1). The conversion holds on the shifted items that decide safety: 27/39 (0.692, 95% CI [0.536, 0.814]), 47/68 (0.691, 95% CI [0.574, 0.788]) and 65/69 (0.942, 95% CI [0.860, 0.977]). When the retry substitutes, it picks the item's paired member on 0.635, 0.725 and 0.665 of those calls (shifted: 0.593, 0.596 and 0.554), so the retried call is usually the near-miss the item was written against, not some other member.

Table 8: **Table 8.** One validation retry on every echoed awkward call.

Model	Retried	Became substitutions	Shifted: became substitutions	Second echo	Clarify / no call
Qwen2.5-7B	80	63 (0.787)	27/39 (0.692)	0	17
Qwen2.5-3B	181	131 (0.724)	47/68 (0.691)	47	2
Mistral-7B	164	155 (0.945)	65/69 (0.942)	6	1

The retry lists the allowed values, which is what a validator's error commonly does, so part of the effect is that the option set is made salient at the moment the model must revise. The direction is nevertheless the deployment-relevant one: the check that blocks a call does not prevent the

substitution, it delays it by one turn, and the retry policy decides whether a loud failure stays visible or becomes a silent substitution.

This qualifies the three-model contrast rather than reversing it. In the no-retry harness the anchor is silent and the other two are loud; under a single retry all three are silent on most of the same items, including the shifted ones. The aggregate safety ranking of a model therefore cannot be read from its no-retry behaviour any more than from its clean accuracy. A pipeline that retries on validation error faces the same silent failure from the 3B and Mistral that it faces from the anchor; a pipeline that does not retry is protected from those two only by refusing to complete the request. The safer retry is one that surfaces the conflict to the caller rather than asking the model to satisfy the schema again.

6. Discussion

6.1 The blind spot is structural, and silence is a policy

The natural reading of the schema-validation result is that the validator should be improved. That reading is wrong: a validator has access only to the call, so it can see that the value is a member but cannot see what the user asked for. No improvement closes the gap, because the information required is not in its input. The literal check has both inputs but compares strings, so it treats a correct paraphrase and a dangerous shift identically. The missing component is a semantic comparison, and §5.9 shows it is not trivial to build.

The three-model result removes the two shortcuts an engineer would like to use. The first is scale: Mistral-7B is the same size as the anchor and resolves the conflict in the opposite way, so a bigger model is not automatically the quieter one and a smaller model is not automatically the loud one. The second is clean-task accuracy: all three models score zero errors on the matched clean items, so an evaluation that screens on ordinary tool-call accuracy ranks them as tied and tells you nothing about this failure. The comprehension control shows how little that tie means: on the same items, the fraction of awkward values a model can map to a member at all ranges from 0.66 to 1.00 while clean accuracy stays at 1.00 for all three.

What actually differs is what each model does when the instruction and the schema disagree — a policy that its authors chose, that is visible nowhere in its accuracy, and that we find varies from "echo the invalid value and be caught" to "rewrite the argument and proceed". For the anchor the policy is the whole story, because the mapping is always available to it; for the two loud models the echo rate also reflects a partial mapping deficit, so their behaviour should not be read as a pure policy contrast (§5.5). An evaluation that measures only schema-validity rates will rank the anchor as the most successful model of the three, at 0.548 substitution against 0.086 and 0.176, because every one of those calls validates and runs. The metric moves the right way while the risk moves the wrong way. This is the sense in which the gate is not merely incomplete but inverting: it rewards exactly the behaviour the caller cannot see. A retry completes the inversion: the same conflict, tried again, moves from a raised error to a valid execution (§5.11).

This is an instance of a broader pattern in machine learning systems, where the failure modes that matter are the ones the monitoring does not represent (Sculley et al., 2015; Breck et al., 2017; Amodei et al., 2016). The contribution here is a concrete, measurable instance in a widely deployed interface.

6.2 Why the comparators fail

The two attempts were designed to be the constructive half of the paper, and they did not work. The reason is instructive. The classification task we posed — is this member an acceptable mapping of this value — is not a similarity task. "Staging but with real data" is lexically and semantically closer to `staging` than to `prod`, yet the correct reading of the user's intent is that they want the data, which is the production characteristic. The decision depends on which dimension of the value the user cares about, and that is not recoverable from the value alone.

An annotation-based evaluation would face the same problem: two annotators can disagree about whether "staging but with real data" maps to `staging` or is a conflict to surface. Our 138/72 split is a single annotator's judgment, and this is the study's main limitation. It does not affect the schema-validation or literal-check results, which depend only on string and membership facts, and it does not affect the substitution rate, which is a count of emitted members. It does affect the shifted detection numbers, which is why we report them under both the full and the unambiguous subset.

6.3 What a pipeline should do

Five practices follow directly from the measurements, none of which requires a new model.

First, do not treat schema validation as a faithfulness check. It answers a question about types, and the paper's central number — 0 of 115 — is what happens when it is asked a question about meaning.

Second, make the conflict visible at the boundary rather than resolving it inside the model. The instruction is the only place where the conflict is still explicit, so a pre-call check that the supplied value is a member or a whitelisted paraphrase turns a silent substitution into a surfaced one, at the cost of a rejection the caller can act on. Being a string check on the *instruction* rather than the call, it runs before the model can paraphrase and does not suffer the false-alarm problem.

Third, be careful how the prompt states the requirement. Telling the model its call must validate, and must never emit a value outside the enum, sounds like a safety instruction; on the anchor it raises silent meaning changes from 0.347 to 0.778. It specifies the *form* of the output and leaves the model to satisfy it, and the cheapest way to satisfy it is to rewrite the argument. A prompt that wants the conflict surfaced has to say so, not merely that the output must be valid. The permission also has to be concrete: adding *"if the value cannot be expressed with the allowed values, do not call the tool; instead ask the user"* cuts the anchor's silent substitutions from 0.548 to 0.348, but the two loud models ignore the same sentence entirely.

Fourth, log the divergence between instruction and call. Even without a classifier, storing both and flagging non-verbatim matches gives an audit trail, and 0.822 false alarms is acceptable for a human-review queue. The distinction between a check that blocks and one that surfaces is what makes the literal check unusable as a gate and useful as an alarm.

Fifth, do not retry a schema-invalid call into compliance. One validation retry converted 0.787, 0.724 and 0.945 of the three models' echoed calls into schema-valid substitutions, and 0.692, 0.691 and 0.942 of the meaning-changing ones. A retry that asks the model to satisfy the schema again is a machine for turning a visible failure into an invisible one. If the pipeline cannot surface the conflict to the caller it should stop; if it can, the caller should choose the member, not the model.

6.4 Limitations

The study uses three models — one family at two scales plus one same-sized model from another family — all in 4-bit quantization on a single laptop GPU; quantization can change behaviour, and we did not run a full-precision control. Three models is enough to show that the silent route is not a scaling law, but not enough to say how common it is; the policy could be rare, or could be typical of models we did not run. The instruction arms use one wording each, measured once; they establish that the route is selectable and that the validity clause alone selects it, not how sensitive that response is to phrasing. The item set is generated rather than drawn from production traffic, so the rates are estimates for a designed distribution, not for any particular deployment. The near-miss labels are single-annotator. The downstream checks are simulated on emitted calls rather than run through a real framework. The comprehension control forces a choice, so it establishes that a mapping is available, not that the model would volunteer one unprompted; the abstention and retry arms use one wording each, and the retry arm replays a single retry rather than a loop and lists the allowed values, as a validator's error commonly does, which makes the option set salient at the moment the model revises. The comparators are trained on 138 positive pairs, which is small, and their failure may be a data-scale result rather than an architectural one; we cannot distinguish those with the data we have. Finally, we measure one conflict type — a value outside a closed enum — and say nothing about numeric ranges, free-text arguments, or multi-field constraints, where the same logic may or may not apply.

7. Conclusion

A schema-valid call is not a faithful call. When an instruction supplies a value a closed enum cannot express, Qwen2.5-7B silently substitutes a near-miss member in 54.8% of cases and changes the meaning in 34.7% of the items where the nearest member is wrong; every such call validates and executes, schema validation detects none, and literal comparison detects most at the price of rejecting four in five correct mappings. The silence is a policy rather than a property of the task. With the schema removed the anchor names an allowed member for every one of the 210 awkward items, so the mapping is available; and merely permitting a clarification, with the instruction to emit a single call unchanged, cuts its silent substitutions to 34.8% and turns 64.8% of the items into questions. A smaller sibling and a same-sized model from another family, both equally perfect on the clean items, echo instead — but their loudness is not a safety property either: one validation retry converts most of their echoes, including most of the meaning-changing ones, into schema-valid substitutions. That is why the risk cannot be screened for by size, by ordinary accuracy, or by whether the base model fails loudly. Two trained semantic comparators fail to supply the missing check at a usable operating point, and the released data makes that negative checkable. The practical conclusion is modest and immediate: the check that gates execution does not measure faithfulness, the conflict should be caught at the instruction before the model can resolve it quietly, and a validation failure should be surfaced to the caller rather than retried into a valid call.

Data and code availability

The fifteen schemas, 420 items, prompts, system-prompt arms, per-item generations, result JSONs, the pre-registration and its follow-up addendum, and every analysis script are released at

`https://github.com/LakshitF/PracticeAI/tree/main/csml/submission`

Every number in §5 is produced by one of those scripts from a released JSON; none is hand-transcribed.

References

Amodei, D., Olah, C., Steinhardt, J., Christiano, P., Schulman, J., & Mané, D. (2016). Concrete problems in AI safety. *arXiv preprint arXiv:1606.06565*.

Bender, E. M., Gebru, T., McMillan-Major, A., & Shmitchell, S. (2021). On the dangers of stochastic parrots: Can language models be too big? *Proceedings of the ACM Conference on Fairness, Accountability, and Transparency*, 610–623.

Beurer-Kellner, L., Fischer, M., & Vechev, M. (2024). Guiding LLMs the right way: Fast, non-invasive constrained generation. *Proceedings of the 41st International Conference on Machine Learning*.

Breck, E., Cai, S., Nielsen, E., Salib, M., & Sculley, D. (2017). The ML test score: A rubric for ML production readiness and technical debt reduction. *IEEE International Conference on Big Data*.

Devlin, J., Chang, M.-W., Lee, K., & Toutanova, K. (2019). BERT: Pre-training of deep bidirectional transformers for language understanding. *Proceedings of NAACL-HLT*, 4171–4186.

Gebru, T., Morgenstern, J., Vecchione, B., Vaughan, J. W., Wallach, H., Daumé III, H., & Crawford, K. (2021). Datasheets for datasets. *Communications of the ACM, 64*(12), 86–92.

Geng, S., Josifoski, M., Peyrard, M., & West, R. (2023). Grammar-constrained decoding for structured NLP tasks without finetuning. *Proceedings of EMNLP*.

Hsieh, C.-Y., Chen, S.-A., Li, C.-L., Fujii, Y., Ratner, A., Krishna, R., Lee, C.-Y., & Pfister, T. (2023). Tool documentation enables zero-shot tool-usage with large language models. *arXiv preprint arXiv:2307.00691*.

Huang, L., Yu, W., Ma, W., Zhong, W., Feng, Z., Wang, H., et al. (2023). A survey on hallucination in large language models. *arXiv preprint arXiv:2311.05232*.

Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., et al. (2023). Survey of hallucination in natural language generation. *ACM Computing Surveys, 55*(12), 1–38.

Jiang, A. Q., Sablayrolles, A., Mensch, A., Bamford, C., Chaplot, D. S., de las Casas, D., et al. (2023). Mistral 7B. *arXiv preprint arXiv:2310.06825*.

Kumar, A. R., & Chopra, P. (2026). Authority bias in language models: Source deference and user agreement are not interchangeable. *Advances in Neural Information Processing Systems*. arXiv:2609.37616.

Kurfali, M., & Östling, R. (2025). Conflicting needles in a haystack: How LLMs behave when faced with contradictory information. *Proceedings of EMNLP*, 34361–34376.

Li, J., Chen, X., Xu, R., et al. (2026). Answer first, evidence second? Uncovering hidden risks in well-structured AI search summaries. *Proceedings of SIGIR*.

Li, M., Zhao, Y., Yu, B., Song, F., Li, H., Guo, H., et al. (2023). API-Bank: A comprehensive benchmark for tool-augmented LLMs. *Proceedings of EMNLP*.

Li, Z., Ye, S., Guo, F., & Dang, Z. (2026). SilentProbe: Measuring silent failure in production APIs used as agent tools. *arXiv preprint arXiv:2609.00035*.

Liang, P., Bommasani, R., Lee, T., Tsipras, D., Soylu, D., Yasunaga, M., et al. (2023). Holistic evaluation of language models. *Transactions on Machine Learning Research*.

Liu, N. F., Lin, K., Hewitt, J., Paranjape, A., Bevilacqua, M., Petroni, F., & Liang, P. (2024). Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics, 12*, 157–173.

Muennighoff, N., Tazi, N., Magne, L., & Reimers, N. (2023). MTEB: Massive text embedding benchmark. *Proceedings of EACL*.

Nogueira, R., & Cho, K. (2019). Passage re-ranking with BERT. *arXiv preprint arXiv:1901.04085*.

Patil, S. G., Zhang, T., Wang, X., & Gonzalez, J. E. (2024). Gorilla: Large language model connected with massive APIs. *Advances in Neural Information Processing Systems*.

Qin, Y., Liang, S., Ye, Y., Zhu, K., Yan, L., Lu, Y., et al. (2024). ToolLLM: Facilitating large language models to master 16000+ real-world APIs. *International Conference on Learning Representations*.

Qwen Team. (2024). Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*.

Ray, J. (2026). The constraint tax: Measuring validity-correctness tradeoffs in structured outputs for small language models. *arXiv preprint arXiv:2605.26128*.

Reimers, N., & Gurevych, I. (2019). Sentence-BERT: Sentence embeddings using Siamese BERT-networks. *Proceedings of EMNLP-IJCNLP*, 3982–3992.

Ruan, Y., Dong, H., Wang, A., Pitis, S., Zhou, Y., Ba, J., et al. (2024). Identifying the risks of LM agents with an LM-emulated sandbox. *International Conference on Learning Representations*.

Schick, T., Dwivedi-Yu, J., Dessì, R., Raileanu, R., Lomeli, M., Zettlemoyer, L., et al. (2023). Toolformer: Language models can teach themselves to use tools. *Advances in Neural Information Processing Systems*.

Scholak, T., Schucher, N., & Bahdanau, D. (2021). PICARD: Parsing incrementally for constrained auto-regressive decoding from language models. *Proceedings of EMNLP*, 9895–9901.

Sculley, D., Holt, G., Golovin, D., Davydov, E., Phillips, T., Dietterich, T., et al. (2015). Hidden technical debt in machine learning systems. *Advances in Neural Information Processing Systems*.

Sigdel, A., & Baral, R. (2026). Schema first tool APIs for LLM agents: A controlled study of tool misuse, recovery, and budgeted performance. *arXiv preprint arXiv:2603.13404*.

Tam, Z. R., Wu, C.-K., Tsai, Y.-L., Lin, C.-Y., Lee, H.-y., & Chen, Y.-N. (2024). Let me speak freely? A study on the impact of format restrictions on performance of large language models. *Proceedings of EMNLP: Industry Track*.

Tan, H., Sun, F., Yang, W., Wang, Y., Cao, Q., & Cheng, X. (2024). Blinded by generated contexts: How language models merge generated and retrieved contexts when knowledge conflicts? *Proceedings of ACL*, 6207–6227.

Thorne, J., Vlachos, A., Christodoulopoulos, C., & Mittal, A. (2018). FEVER: A large-scale dataset for fact extraction and verification. *Proceedings of NAACL-HLT*, 809–819.

Usman, R. M. (2026). PhantomFill: When the form demands an answer, language models invent one. *arXiv preprint arXiv:2607.20492*.

Wei, J., Wang, X., Schuurmans, D., Bosma, M., Ichter, B., Xia, F., et al. (2022). Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems*.

Willard, B. T., & Louf, R. (2023). Efficient guided generation for large language models. *arXiv preprint arXiv:2307.09702*.

Yang, H., Song, W., Kim, T., Song, J., Park, S., & Jo, Y. (2026). Agents' overreliance on unreliable tools. *arXiv preprint arXiv:2609.05587*.

Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., & Cao, Y. (2023). ReAct: Synergizing reasoning and acting in language models. *International Conference on Learning Representations*.

Zheng, C., Zhou, H., Meng, F., Zhou, J., & Huang, M. (2023). Large language models are not robust multiple choice selectors (PriDe). *arXiv preprint arXiv:2309.03882*.

Zheng, L., Chiang, W.-L., Sheng, Y., Zhuang, S., Wu, Z., Zhuang, Y., et al. (2023). Judging LLM-as-a-judge with MT-Bench and Chatbot Arena. *Advances in Neural Information Processing Systems*.

Appendix A. Use of AI in this project

Claude (Anthropic) was used as a research assistant throughout. Its uses were: literature orientation and novelty checking (locating the closest prior work and verifying that the enum-conflict cell was unoccupied); writing the experiment harness and analysis scripts; running the analyses; and drafting and editing this manuscript. The author designed the study, wrote the item set and schemas, made the labelling decisions, specified the pre-registered gates, and verified every reported number against the released JSON artifacts. All model-generated text in the pipeline — the tool calls the paper studies — was produced by the local Qwen and Mistral models, not by Claude. No number in this paper was written by a language model without being checked against a result file.